



3099704: AI for Digital Health



Classification with Tree-Based Models

Prof. Peerapon Vateekul, Ph.D.

Peerapon.v@chula.ac.th



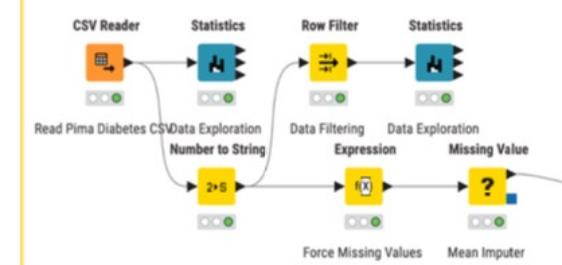
Outlines

- Supervised Learning
- Decision Tree
- Random Forest
- Classification Performance

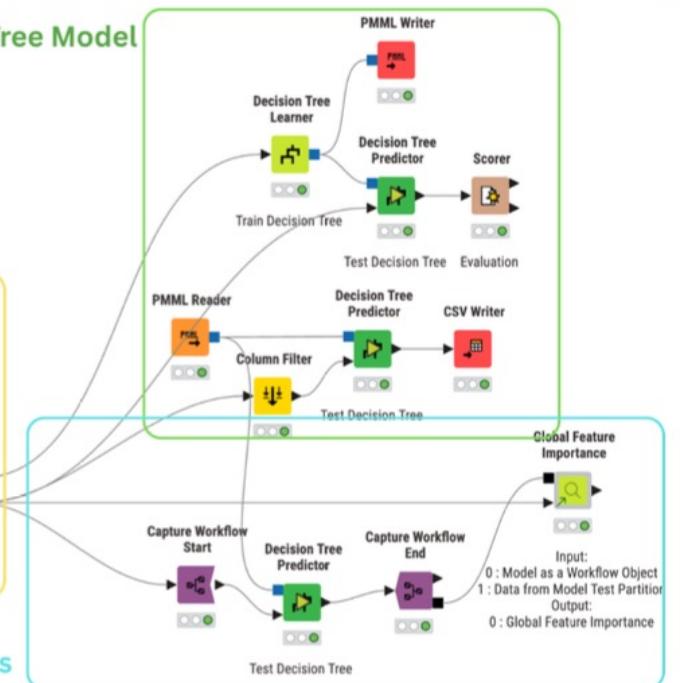
KNIME Instructions

2) Decision Tree Model

1) Data Preparation



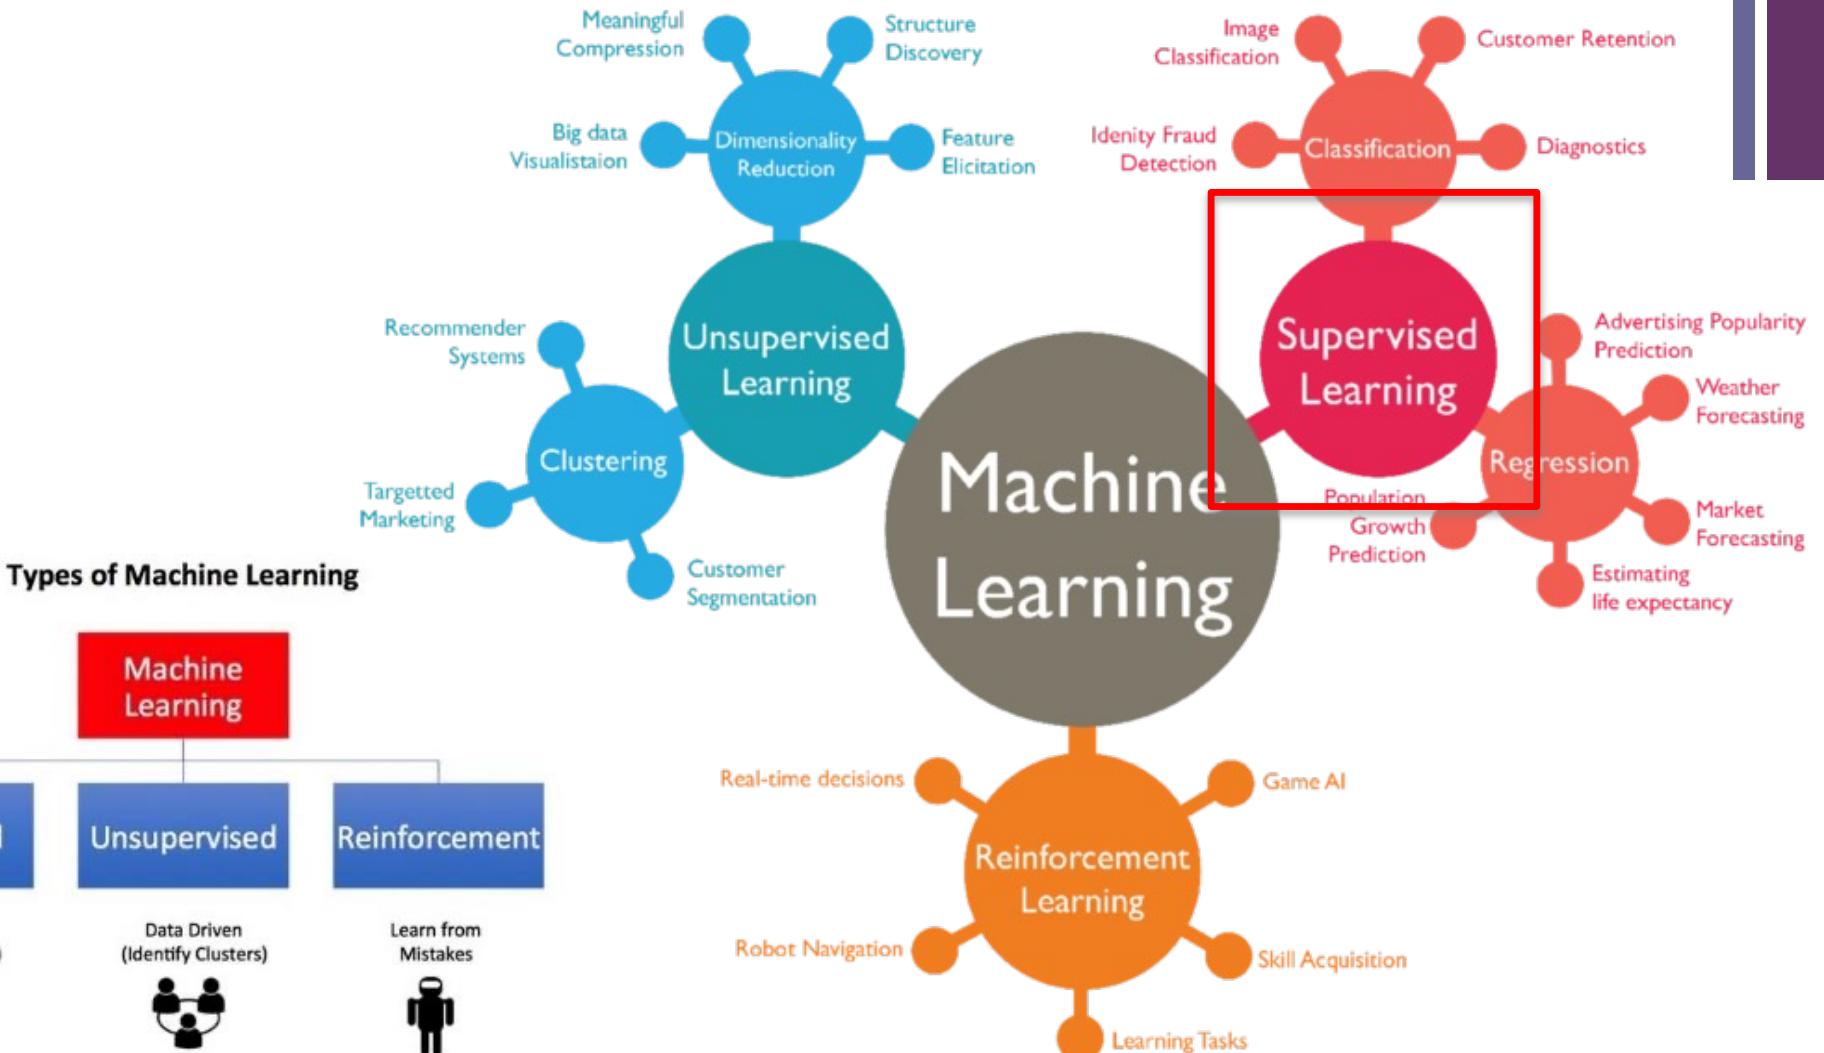
3) Decision Tree Model and Feature Importances





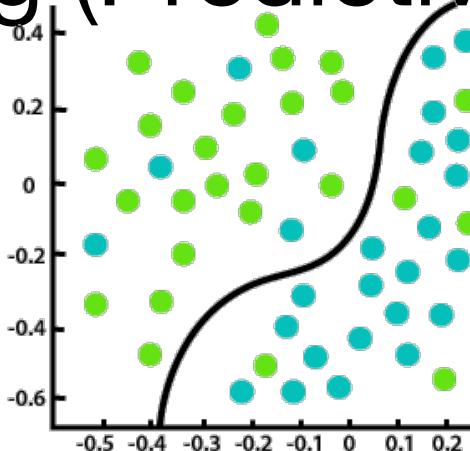
Supervised Learning (Predictive Task)

+ Machine Learning



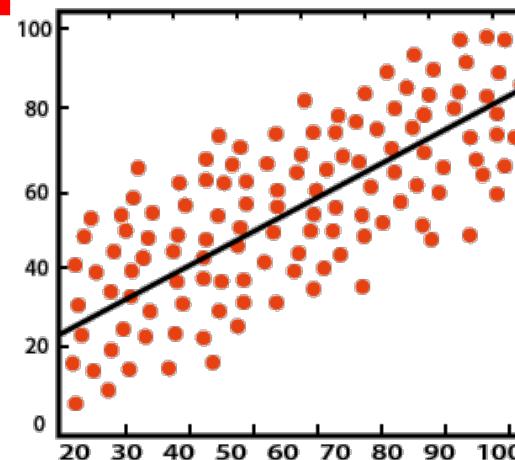
Supervised Learning (Predictive Task)

inputs					target
Age	Temp	Gender	Smell	Covid	
25	39.0	Female	No	Yes	
35	38.9	Female	No	Yes	
32	36.5	Male	Yes	No	



- Target is **categorical** variable.
- Example
- Covid diagnosis (yes/no)
- Disease diagnosis from gait information:
 - 1) Normal,
 - 2) Sick/Knee OA
 - 3) Sick/Parkinson

Classification



- Goal: To learn **a prediction model** mapping from inputs to output.
- Data without label (answer) is meaningless!
- Label should be provided by experts!

- Target is **numeric** variable.
- Example
- **PD's state** diagnosis from movement data.
- **Glucose level** prediction from breath particles.

Regression



There are two main processes: Train/Test

1) Training Phase: Model Construction

Training Data



Age	Income	inputs		Purchase	target
		Gender	Province		
25	25,000	Female	Bangkok	Yes	
35	50,000	Female	Nontaburi	Yes	
32	35,000	Male	Bangkok	No	

2) Testing Phase: Model Evaluation, Model Assessment

Also called “prediction, inference, scoring”

Testing Data



Age	Income	Gender	Province	Purchase
25	25,000	Female	Bangkok	?



Prediction Algorithms

- Decision Tree
- (Logistic) Regression
- Neural Networks (NN)
- kNN
- Support Vector Machine
- Deep Learning

BASIC REGRESSION

- **LINEAR** linear_model.LinearRegression()
Lots of numerical data
- **LOGISTIC** linear_model.LogisticRegression()
Target variable is categorical

CLASSIFICATION

- **NEURAL NET** neural_network.MLPClassifier()
Complex relationships. Prone to overfitting
Basically magic.
- **K-NN** neighbors.KNeighborsClassifier()
Group membership based on proximity
- **DECISION TREE** tree.DecisionTreeClassifier()
If/then/else. Non-contiguous data
Can also be regression
- **RANDOM FOREST** ensemble.RandomForestClassifier()
Find best split randomly
Can also be regression
- **SVM** svm.SVC() svm.LinearSVC()
Maximum margin classifier. Fundamental Data Science algorithm
- **NAIVE BAYES** GaussianNB() MultinomialNB() BernoulliNB()
Updating knowledge step by step with new info



Scikit-learn: Machine learning library in Python

- Provides many machine learning tools with a common **Estimator interface**
- Built in helpers for common **ML tasks** (e.g., metrics, preprocessing)
- Easily combine algorithms to make **a complex pipeline**
- Relies heavily on **numpy** and **scipy**, often used with **pandas**



How do you pronounce the project name?

sy-kit learn. sci stands for science!

Why scikit?

There are multiple scikits, which are scientific toolboxes built around SciPy. You can find a list at <https://scikit.appspot.com/scikits>. Apart from scikit-learn, another popular one is **scikit-image**.

Classification

Identifying to which category an object belongs to.

Applications: Spam detection, Image recognition.

Algorithms: SVM, nearest neighbors, random forest, ...

— Examples

Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, Stock prices.

Algorithms: SVR, ridge regression, Lasso, ...

— Examples

Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation, Grouping experiment outcomes

Algorithms: k-Means, spectral clustering, mean-shift, ...

— Examples

Dimensionality reduction

Reducing the number of random variables to consider.

Applications: Visualization, Increased efficiency

Algorithms: PCA, feature selection, non-negative matrix factorization.

— Examples

Model selection

Comparing, validating and choosing parameters and models.

Goal: Improved accuracy via parameter tuning

Modules: grid search, cross validation, metrics.

— Examples

Preprocessing

Feature extraction and normalization.

Application: Transforming input data such as text for use with machine learning algorithms.

Modules: preprocessing, feature extraction.

— Examples

<http://scikit-learn.org/stable/index.html>



Estimator Interface

Decision Trees

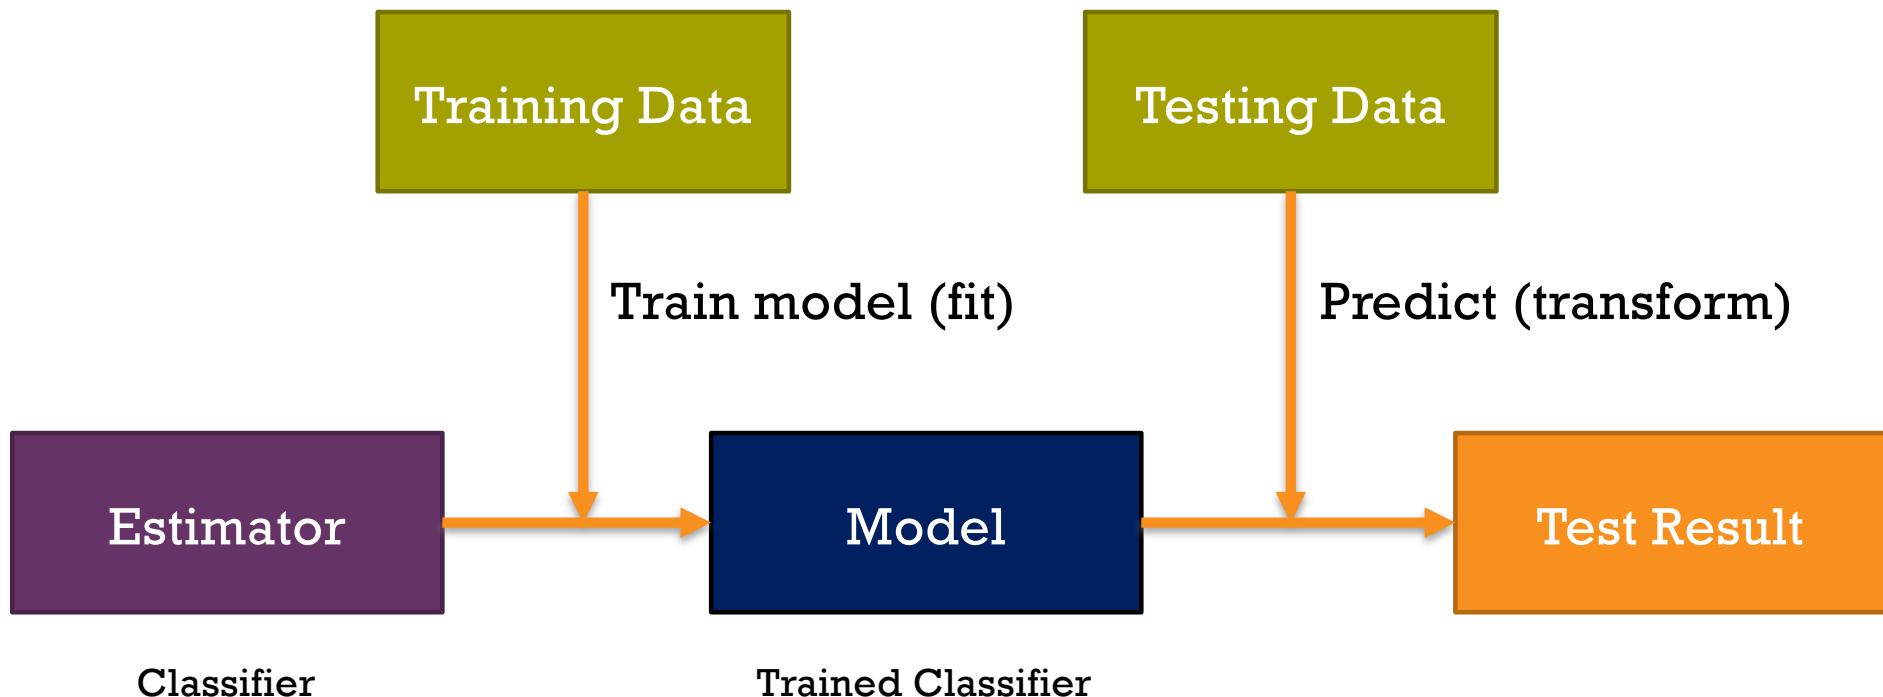
We'll start just by training a single decision tree.

```
In [8]: from sklearn.tree import DecisionTreeClassifier
In [9]: dtree = DecisionTreeClassifier(min_samples_leaf=10, criterion='entropy')
In [10]: dtree.fit(X_train,y_train)
Out[10]: DecisionTreeClassifier(class_weight=None, criterion='entropy', max_depth=None,
                                 max_features=None, max_leaf_nodes=None,
                                 min_impurity_decrease=0.0, min_impurity_split=None,
                                 min_samples_leaf=10, min_samples_split=2,
                                 min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                                 splitter='best')
```

Prediction and Evaluation

Let's evaluate our decision tree.

```
[11]: predictions = dtree.predict(X_test)
[12]: from sklearn.metrics import classification_report,confusion_matrix
[13]: print(classification_report(y_test,predictions))
      precision    recall  f1-score   support
absent       0.85     0.85     0.85      20
present      0.40     0.40     0.40       5
avg / total   0.76     0.76     0.76      25
```

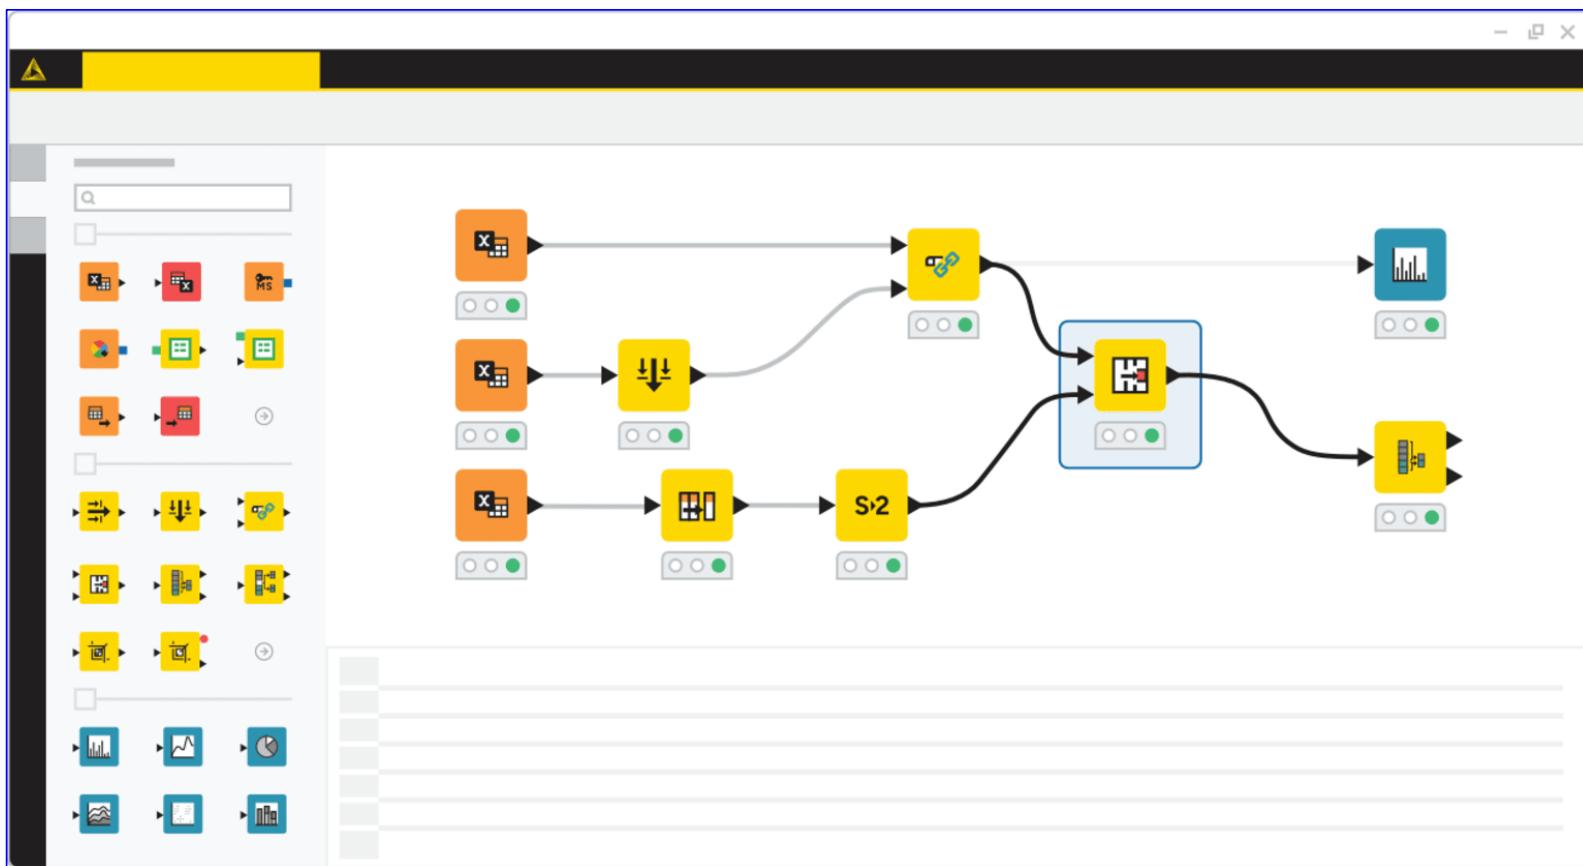




Low-Code/No-Code Software



10

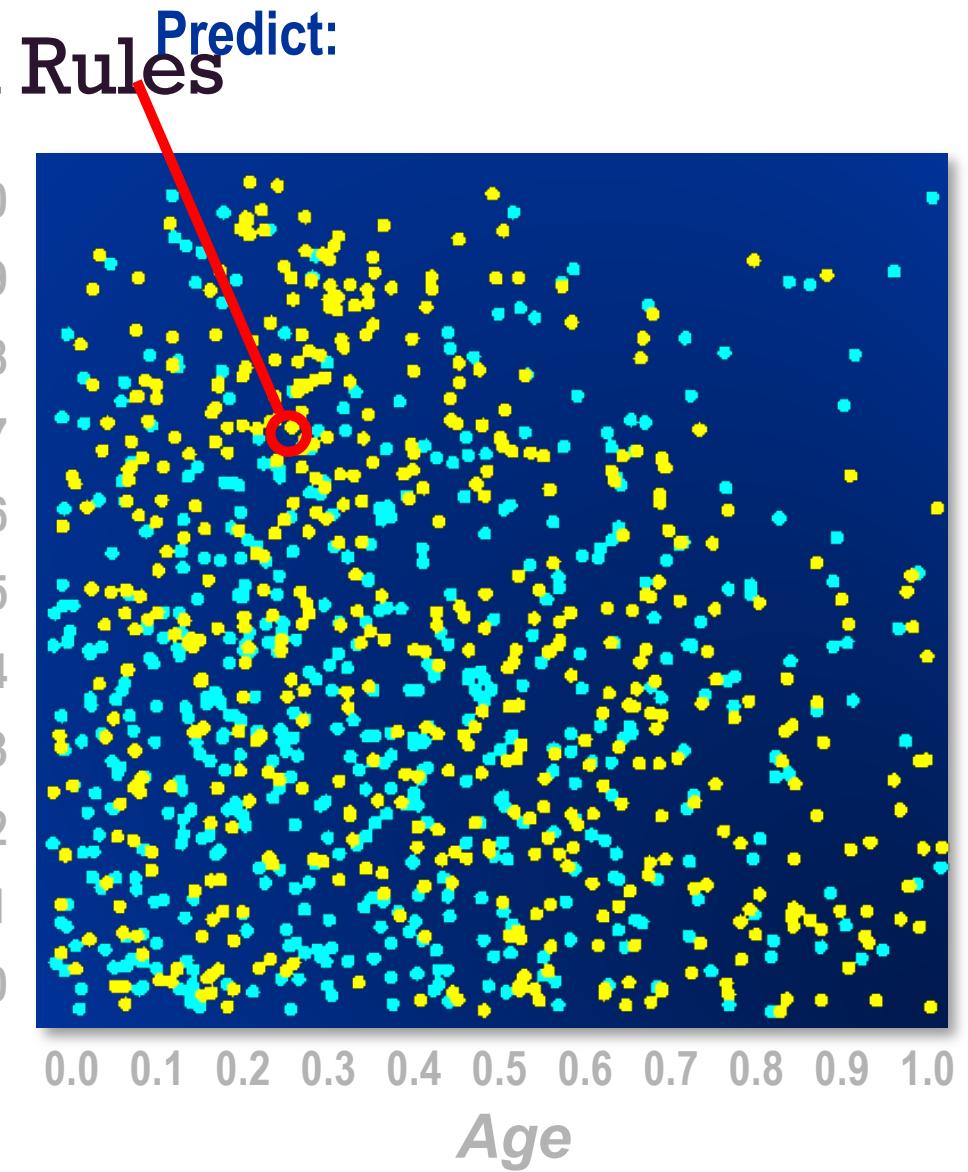
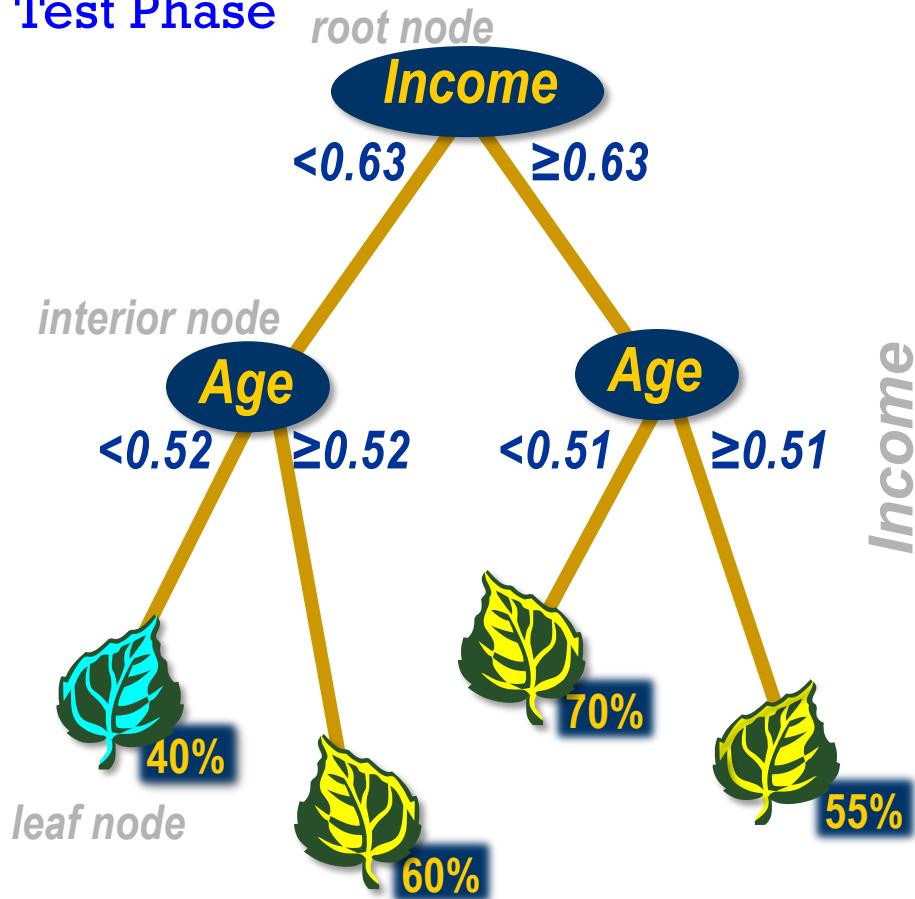


+

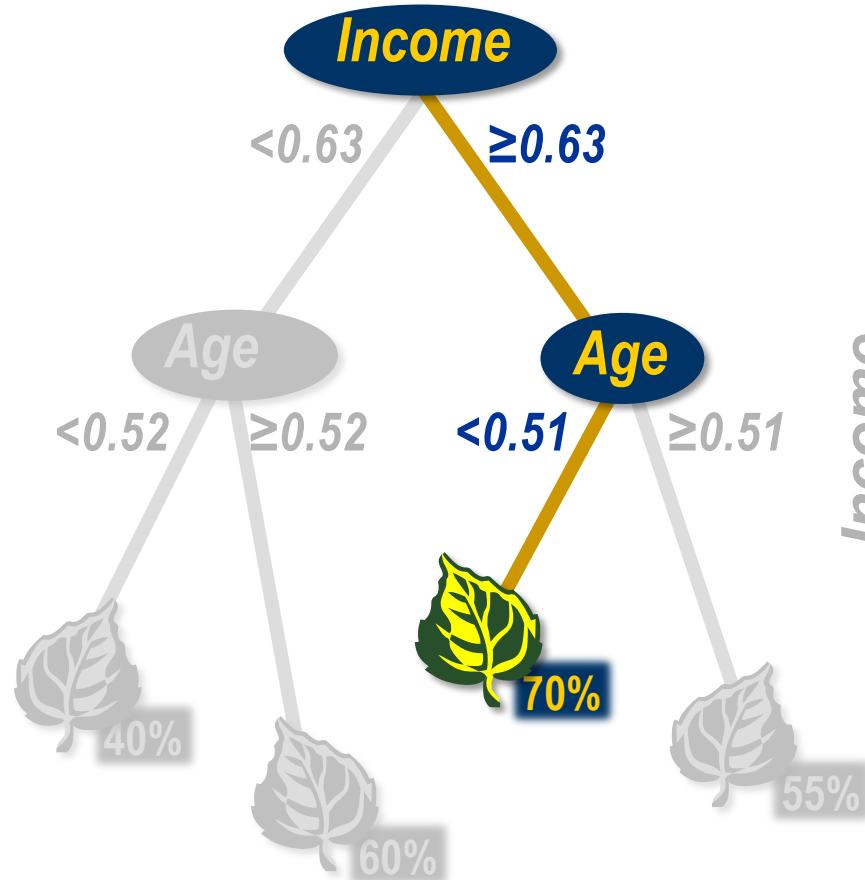
Decision Tree

1) Decision Tree Prediction Rules

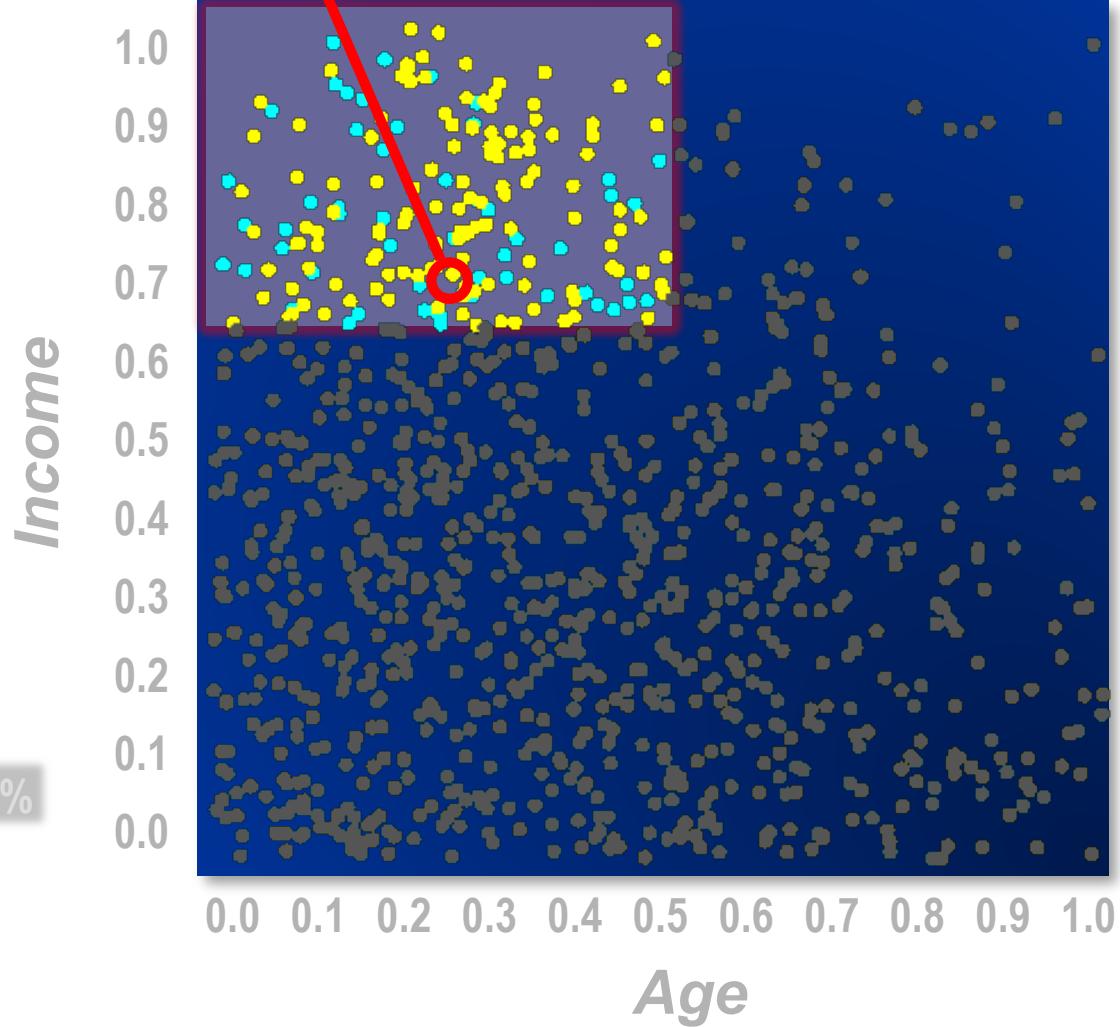
Test Phase



Decision Tree Prediction Rules



Predict: Decision = ●
Estimate = 0.70



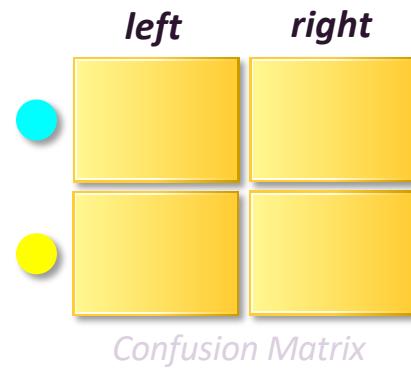
Model Essentials: Decision Trees

- ▶ Predict cases.
- ▶ Select useful predictors.

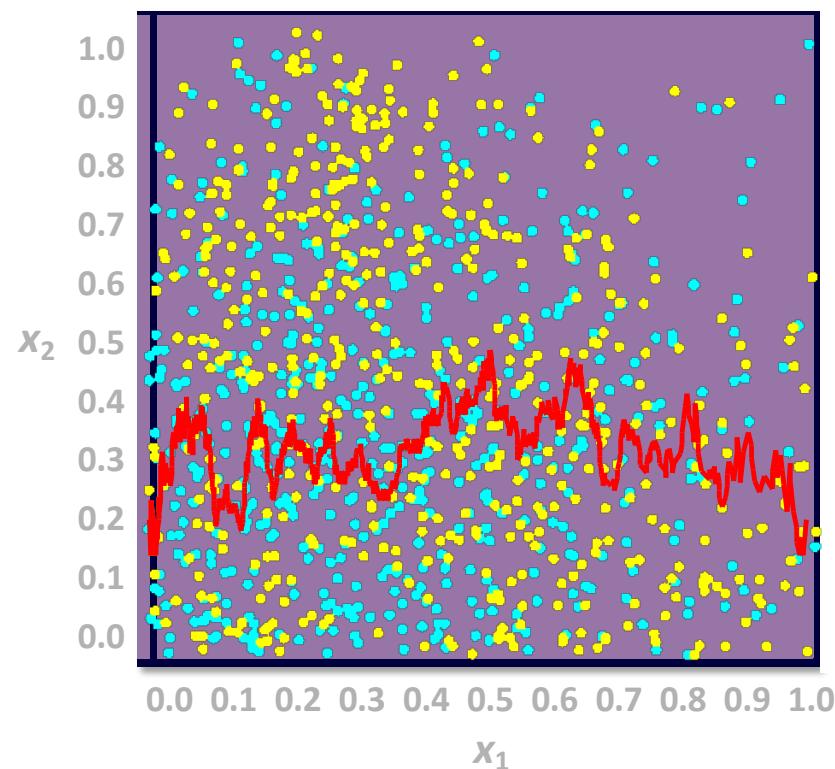
Prediction rules

Split search

Decision Tree Split Search

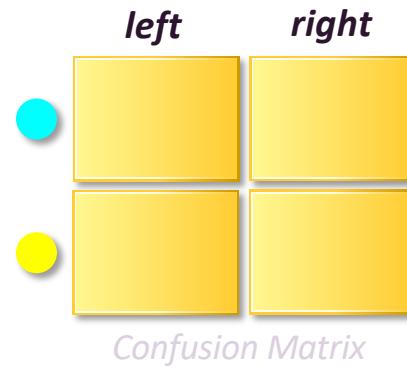


**Calculate information
gain on partitions
on input x_1 .**

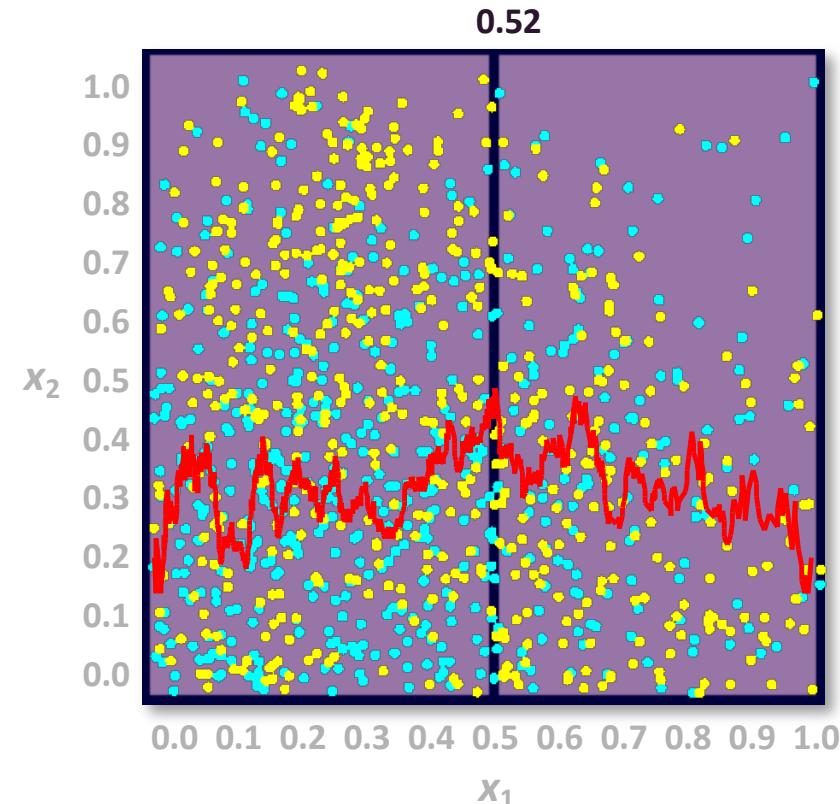


...

Decision Tree Split Search



**Calculate the gain
of every partition
on input x_1 .**



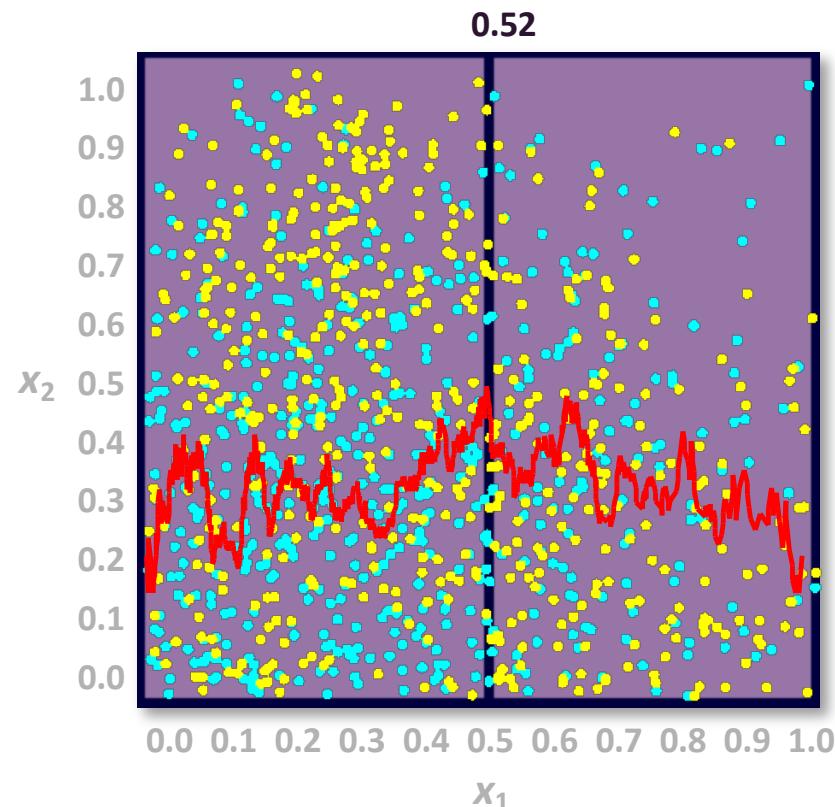
...

Decision Tree Split Search

	<i>left</i>	<i>right</i>	
	53%	42%	
	47%	58%	

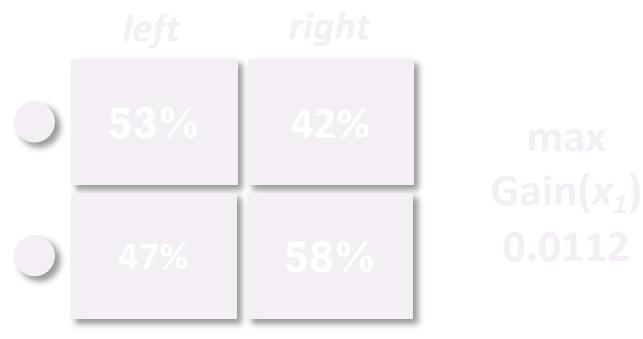
**max
 $\text{gain}(x_1)$
0.0112**

Select the partition with the maximum gain.

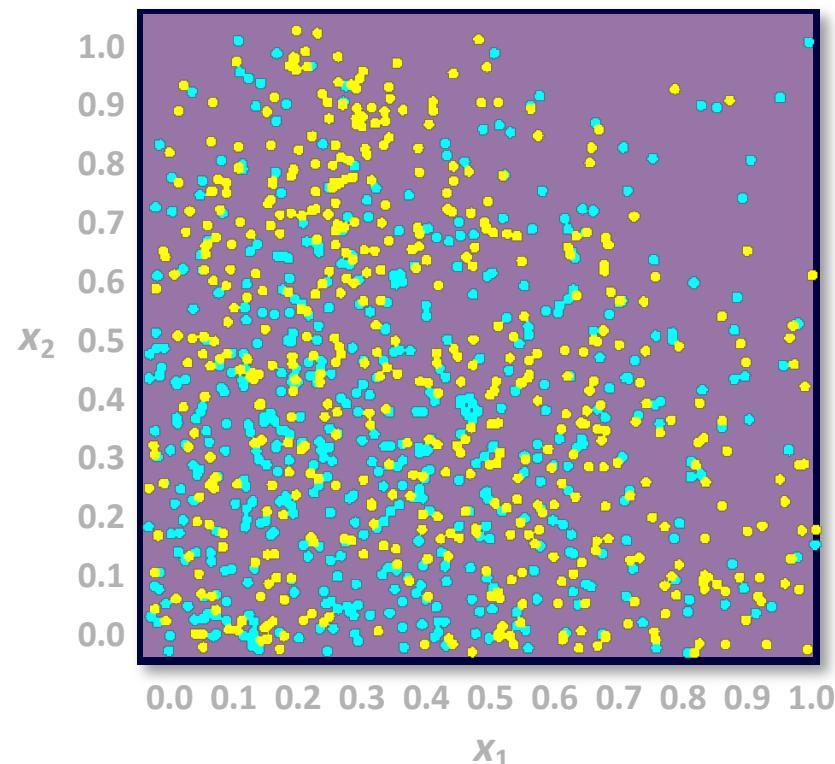


...

Decision Tree Split Search

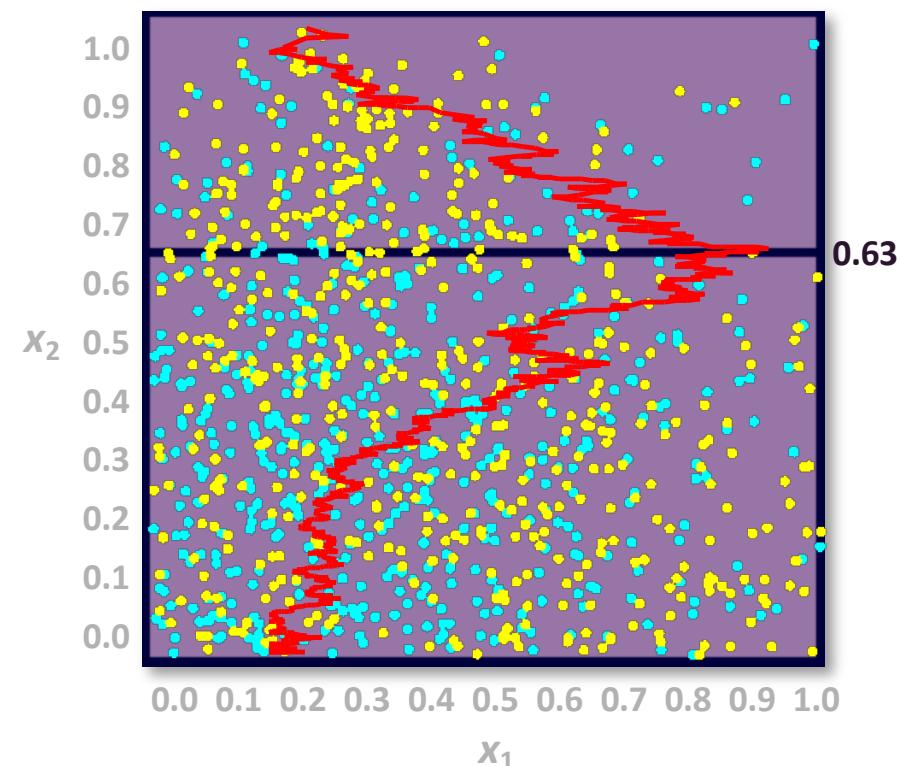
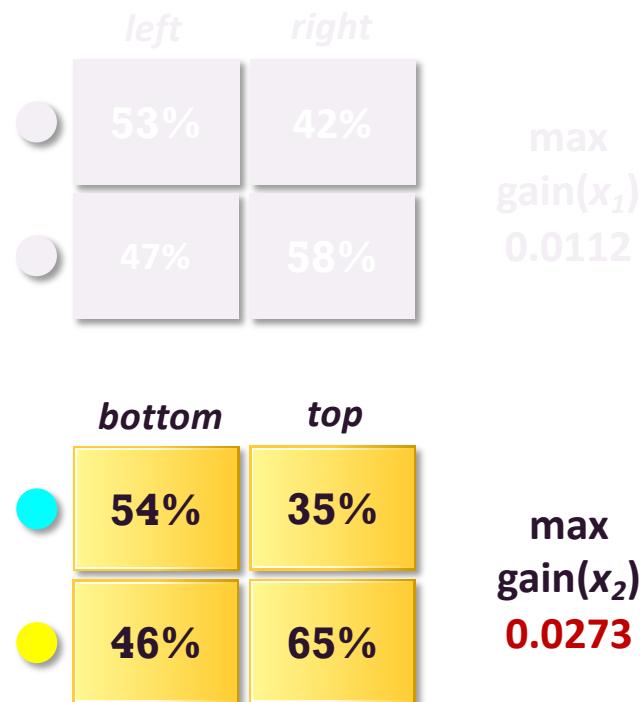


Repeat for input x_2 .



...

Decision Tree Split Search

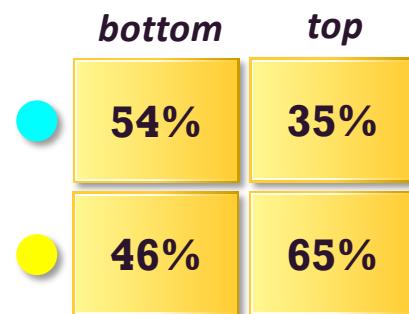


...

Decision Tree Split Search



**max
 $gain(x_1)$**
0.0112

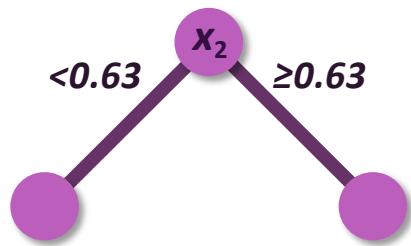


**max
 $gain(x_2)$**
0.0273

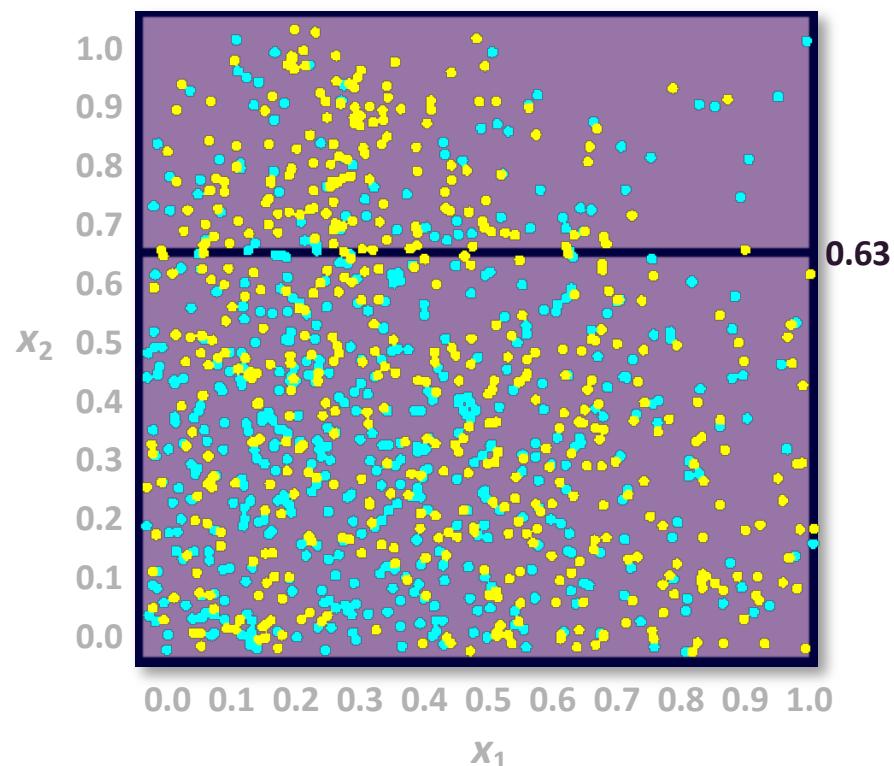


...

Decision Tree Split Search

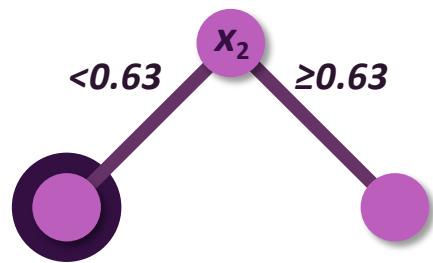


**Create a partition rule
from the best partition
across
all inputs.**

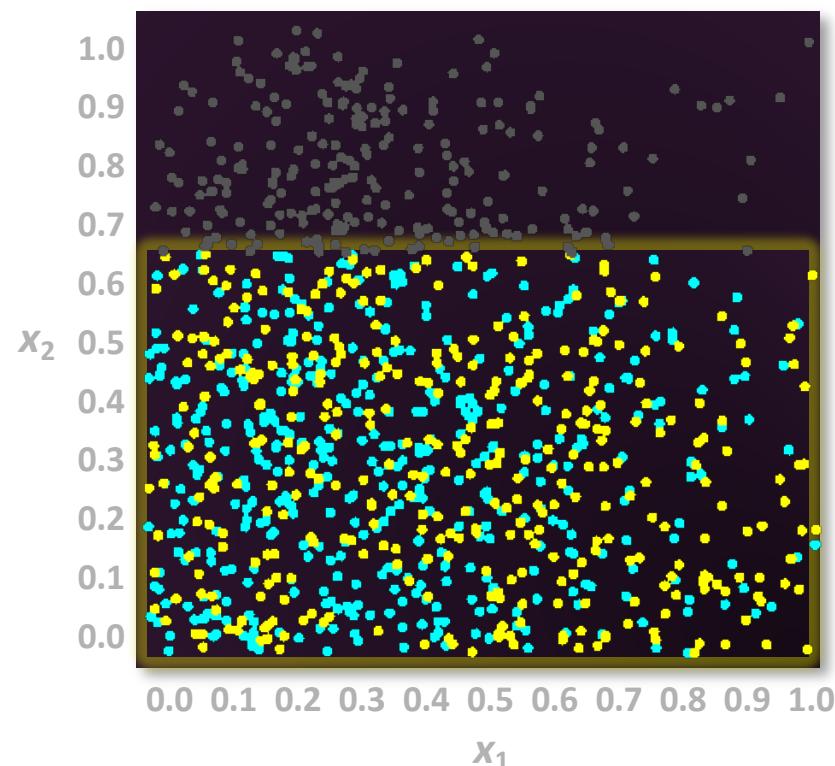


...

Decision Tree Split Search

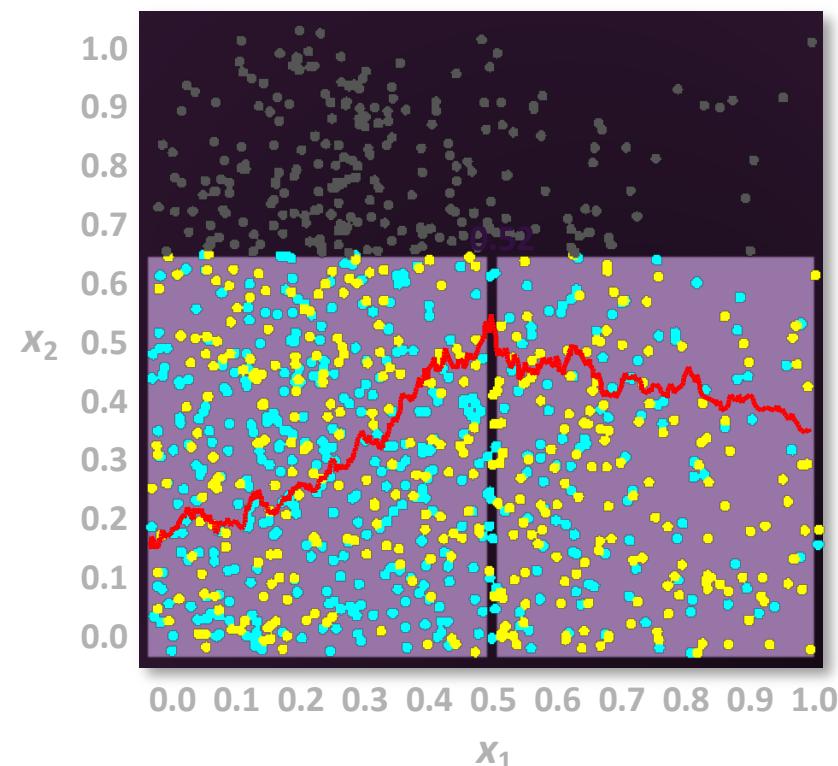


**Repeat the process
in each subset.**



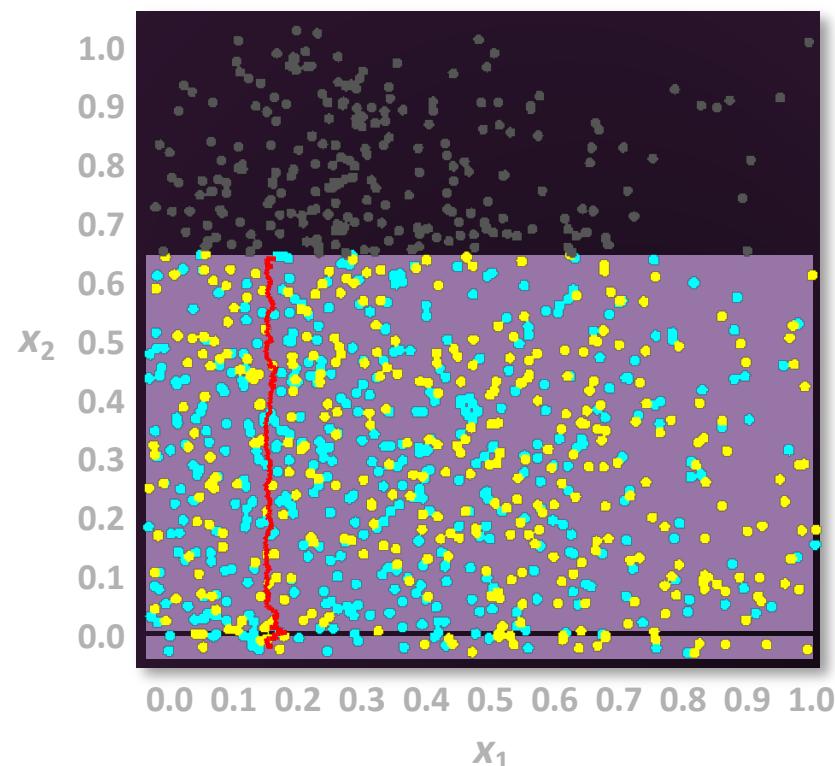
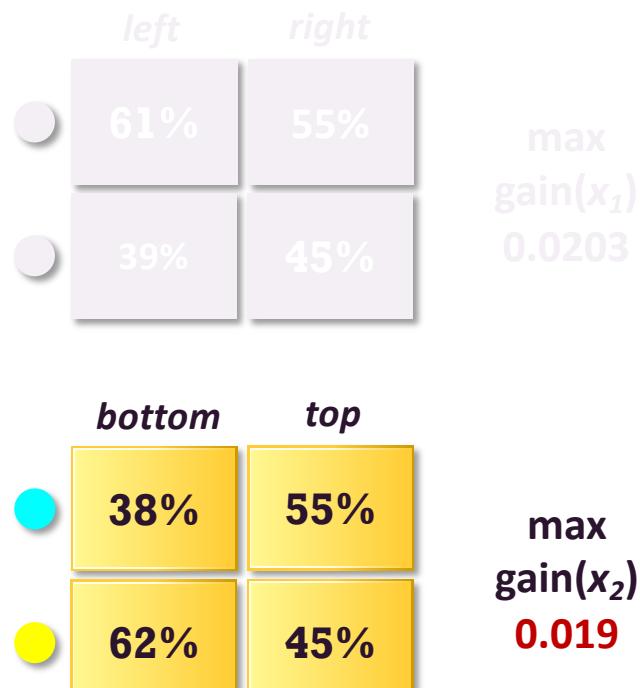
...

Decision Tree Split Search



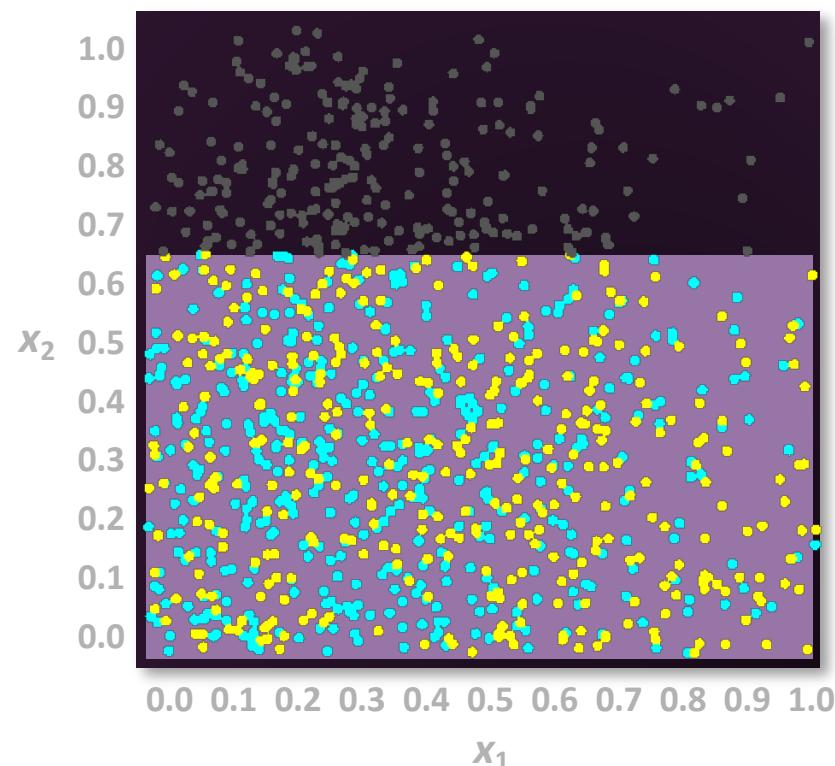
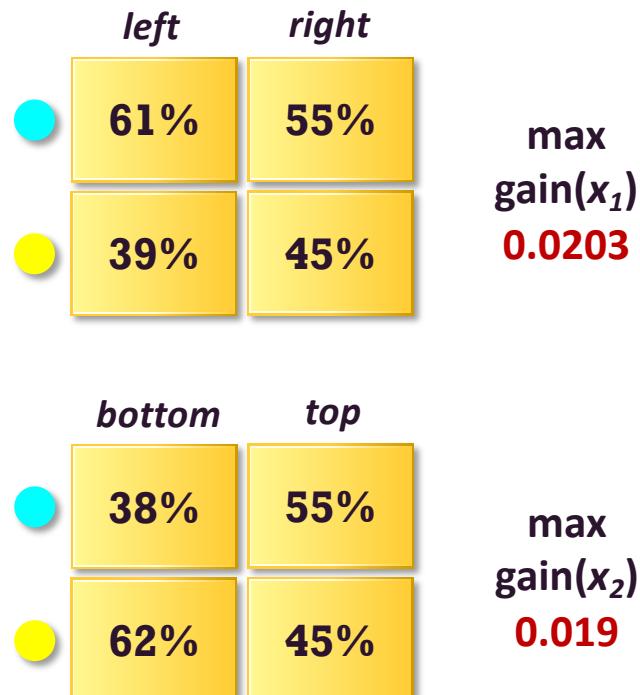
...

Decision Tree Split Search



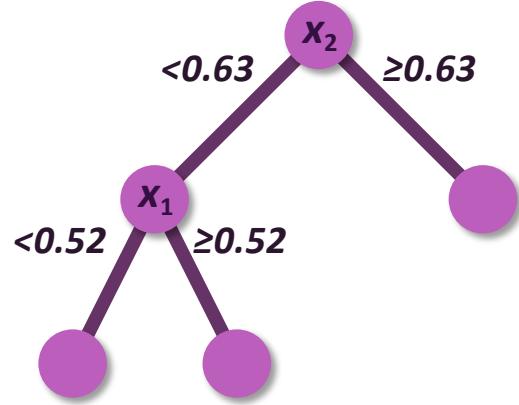
...

Decision Tree Split Search

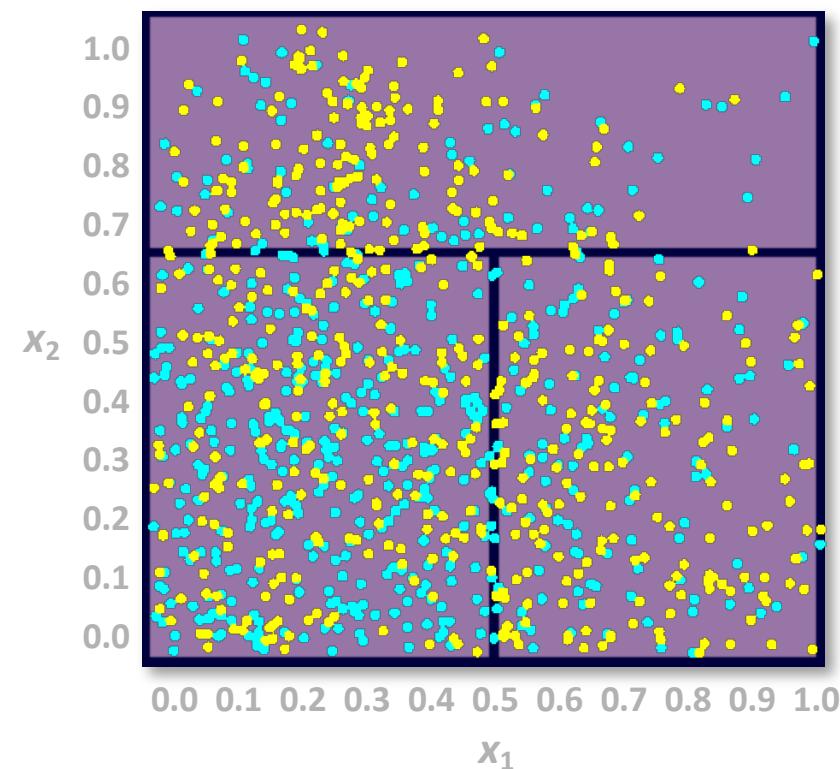


...

Decision Tree Split Search

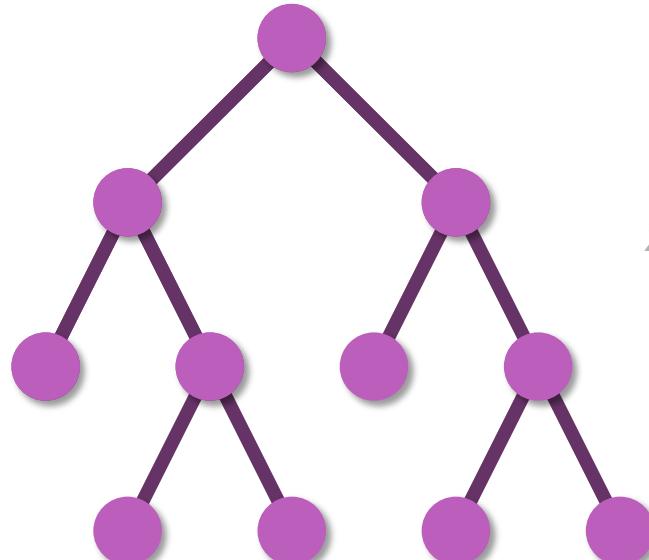


Create a second
partition rule.

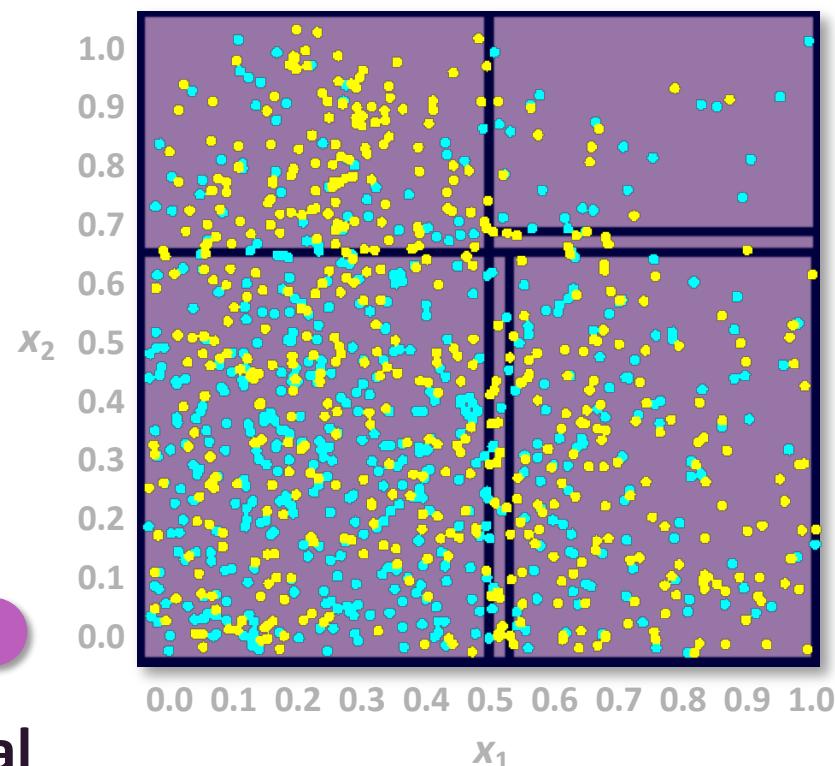


...

Decision Tree Split Search



Repeat to form a maximal tree.





1) Entropy (impurity)

- **Entropy** is a measure of disorder or uncertainty and the goal of machine learning models and general is to reduce uncertainty.

$$\text{Entropy} = \sum_{i=1}^n -p_i \log_2 p_i$$

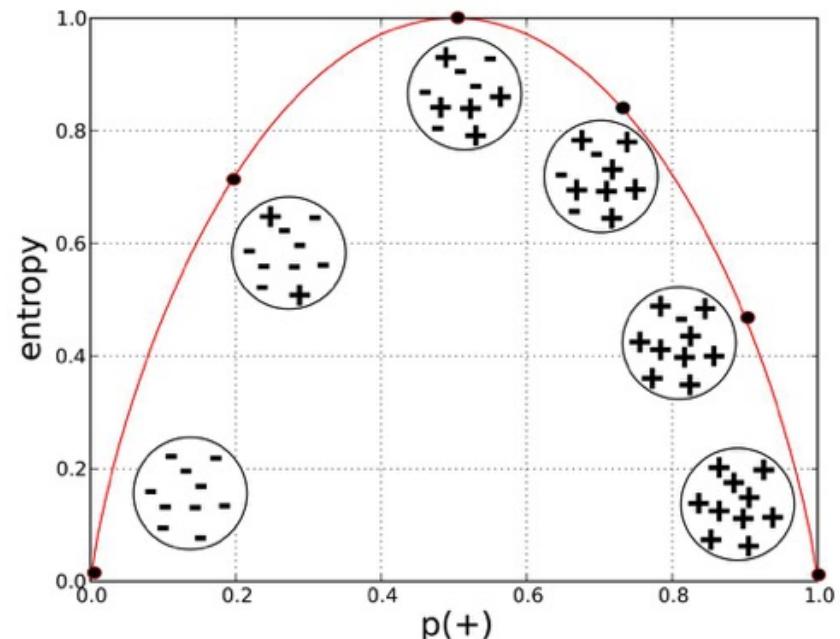
`scipy.stats.entropy`

`scipy.stats.entropy(pk, qk=None, base=None, axis=0)`
Calculate the entropy of a distribution for given probability values.

If only probabilities `pk` are given, the entropy is calculated as `S = -sum(pk * log(pk), axis=axis)`.

If `qk` is not `None`, then compute the Kullback-Leibler divergence `S = sum(pk * log(pk / qk), axis=axis)`.

This routine will normalize `pk` and `qk` if they don't sum to 1.



Source: Data Science for Business: What You Need to Know about Data Mining and Data-Analytic Thinking



Information Gain

Before - After

- Which one is Better ?

Split w/ Age: 70
 \sum Entropy: 0.350

Split w/ Age: 50
 \sum Entropy: 0.348

- Information Gain: measure the reduction of this disorder in our target variable/class given additional information

$$\text{InformationGain} = \text{Entropy}(\text{before}) - \sum \text{Entropy}(\text{after})$$

- “Before” = Entropy of Parent Node
“After” = Entropy of Child Nodes



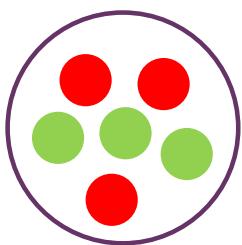
2) Gini Impurity

Gini Reduction = Before - After

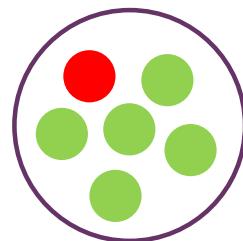
- Another way to measure how well a splitting feature.

$$Gini = 1 - \sum_{i=1}^n (P_i)^2$$

When P is the probability of class i in data-set.



$$\begin{aligned} Gini &= 1 - ((3/6)^2 + (3/6)^2) \\ &= 0.5 \end{aligned}$$



$$\begin{aligned} Gini &= 1 - ((5/6)^2 + (1/6)^2) \\ &= 0.28 \end{aligned}$$

- Easy to calculation, may take less time to build in large dataset.

Source: <https://towardsdatascience.com/understanding-decision-tree-classification-with-scikit-learn-2ddf272731bd>

Types of Decision Tree

Algorithm	Splitting Measure
ID3	Entropy
C4.5	Gain Ratio
CART	Gini index
CHAID	Chi-squared test

Important Parameters in Decision Tree

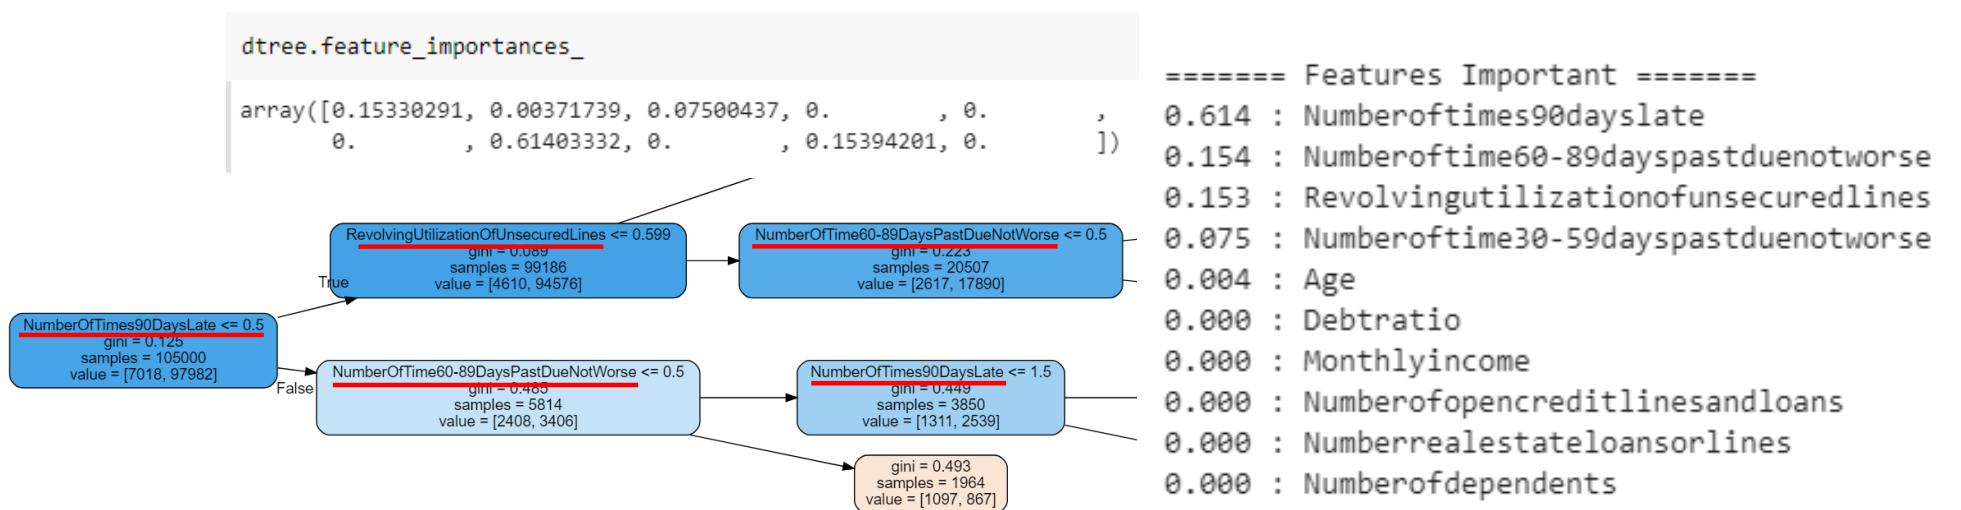


- Splitting measure (criterion) : gini / entropy
- Maximum depth : ~5-10 (depend on number of feature)
- Maximum leaf nodes : depend on number of class (target) and feature
- Minimum sample split : 5 – 20% (depend on number of data)

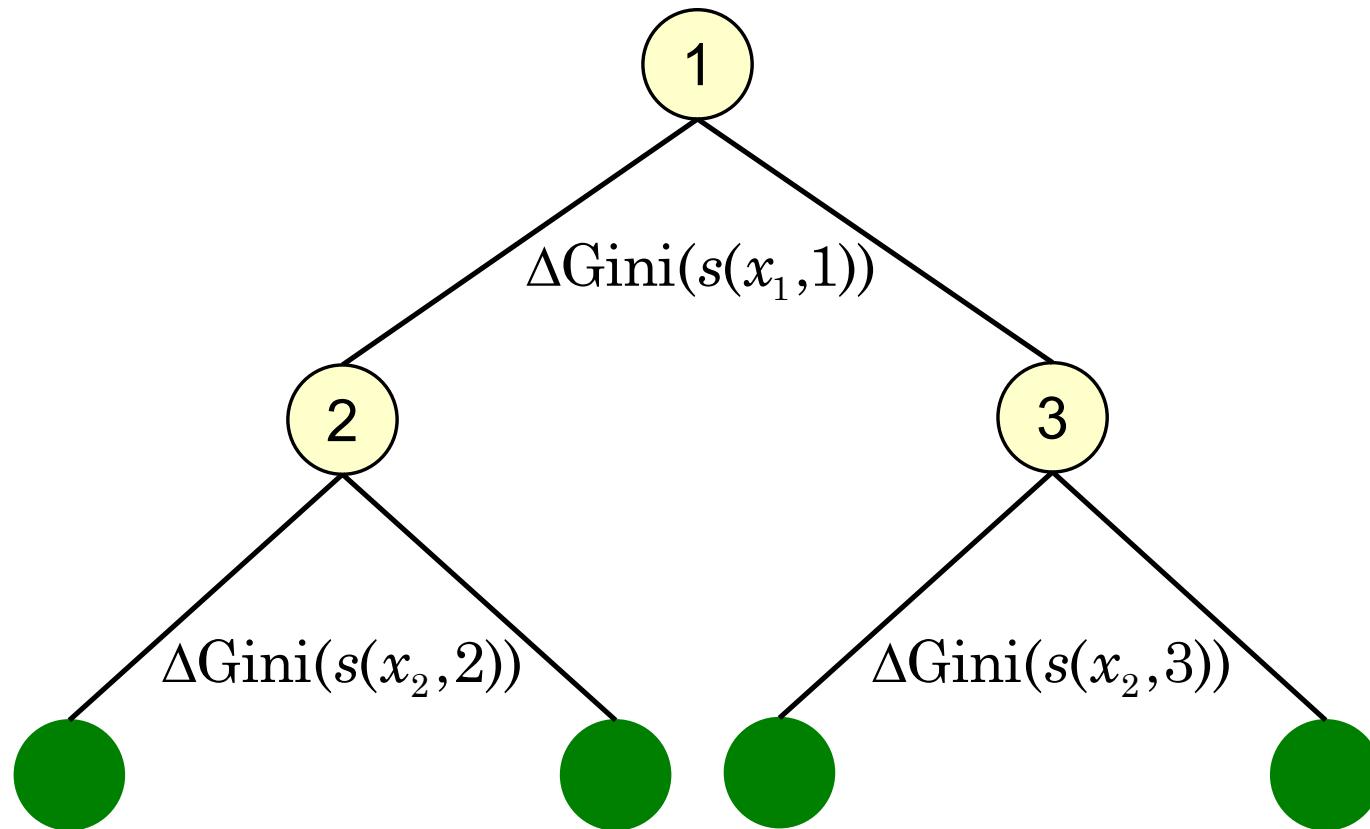
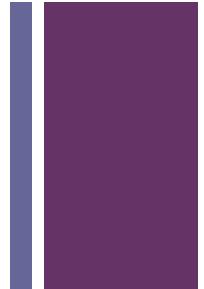


Which features are important ?

- Check how important by `.feature_importances_`
- The importance of a feature is computed as the (normalized) total reduction of the criterion brought by that feature. It is also known as **the gini importance**.



+ Variable Importance



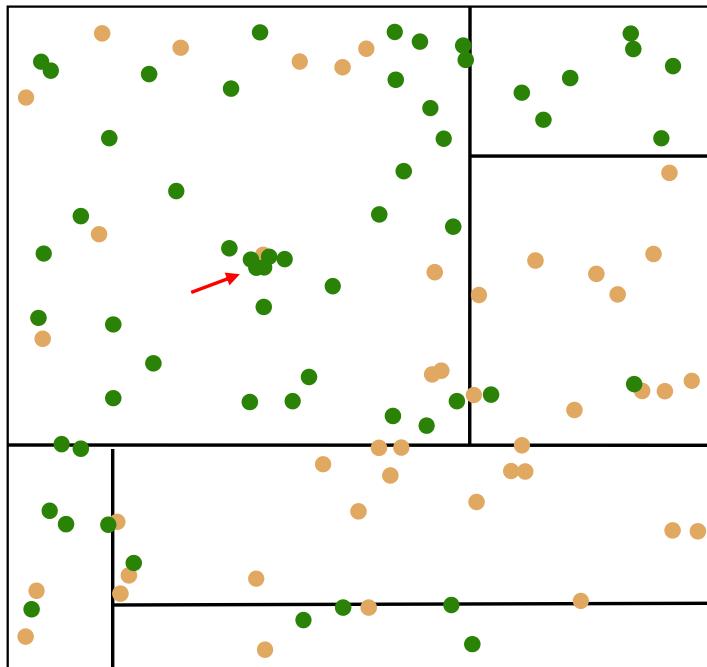
+

Random Forest

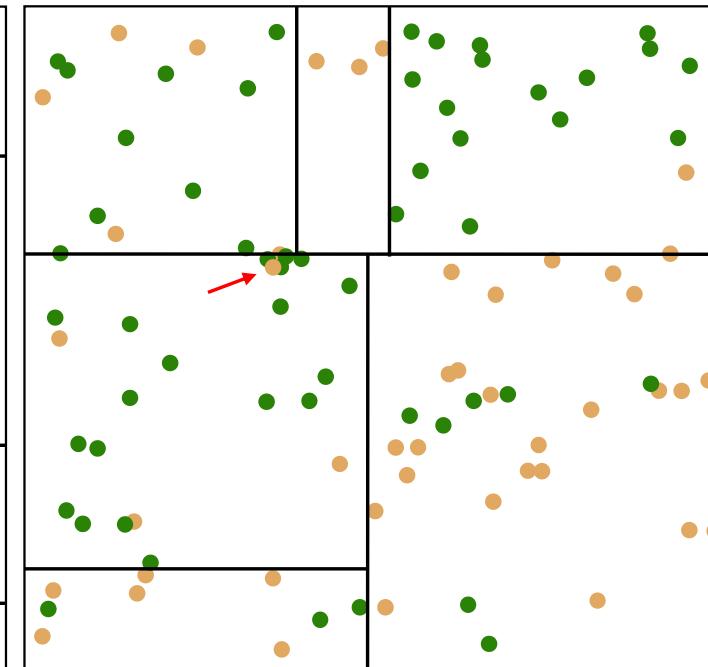
+ Instability

36

One reversal



Accuracy = 81%

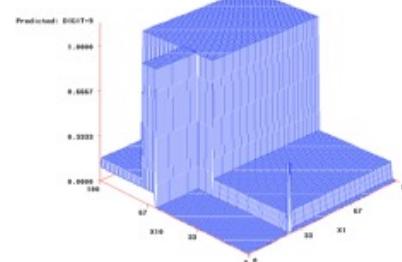


Accuracy = 80%

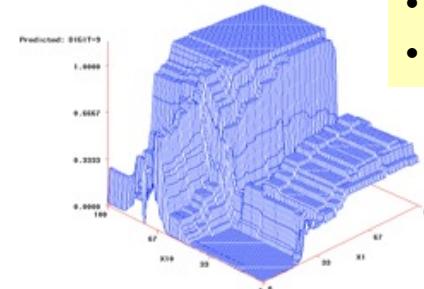
...

Random Forest Algorithm (RF)

- Forest algorithm samples the **rows** and the **columns** at each step.
- Forest takes bootstrap samples of the **rows** in training data (sampling with replacement)
- At each step, a set of variables (**columns**) is sampled.
- This increases variation among trees in the **ensemble** often leads to improved prediction accuracy.



Single Tree

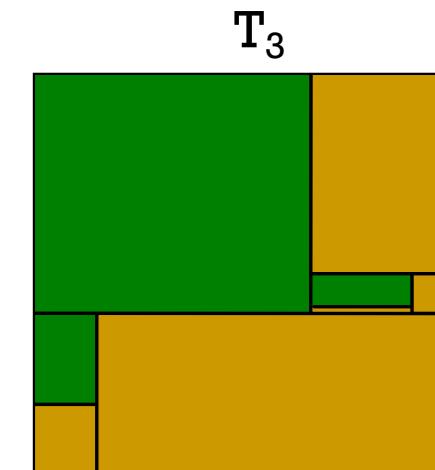
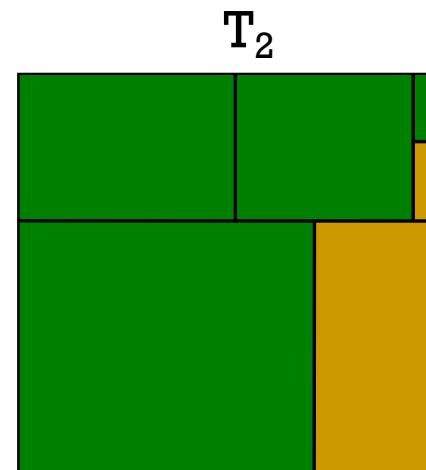
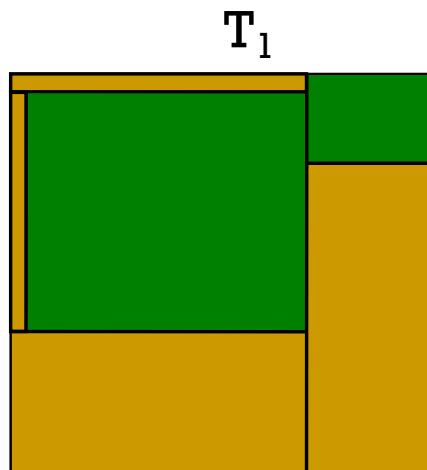


Random Forest

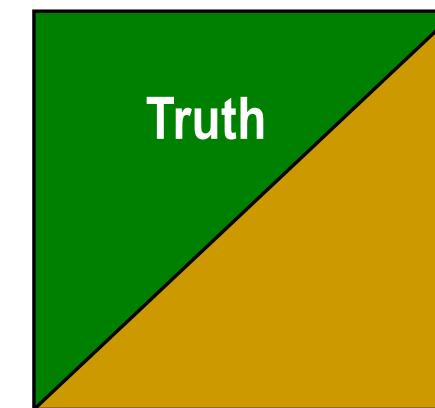
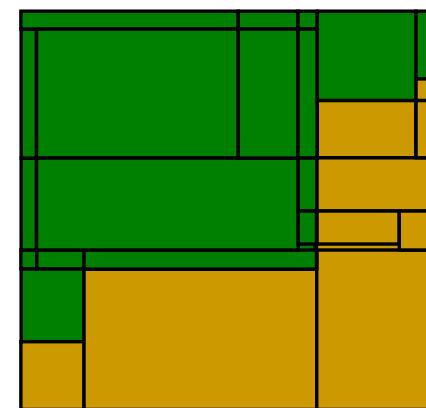
Important Params:

- #Tree
- #Columns (variables)
- #Rows (examples)

Combine



$$\text{avg}(T_1, T_2, T_3) =$$



...



Classification Performance



Evaluation (Train/Test Split)

Training Data

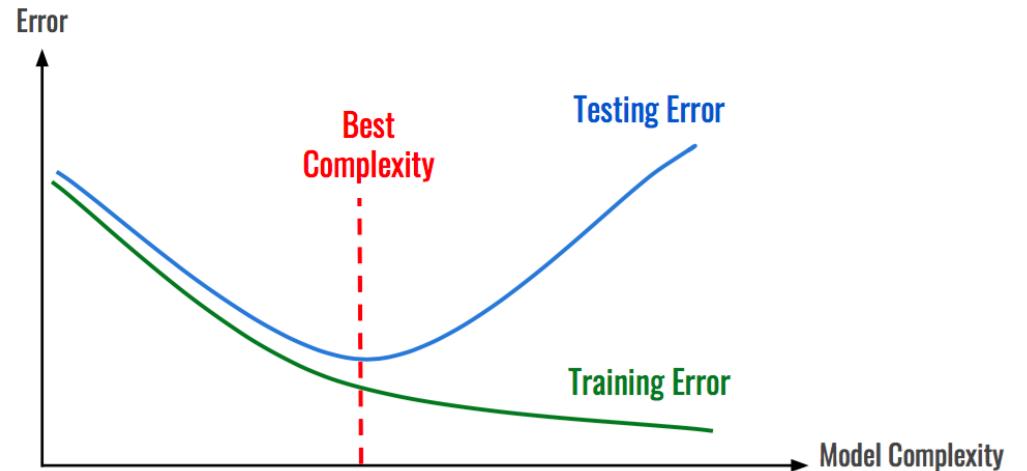


Age	Income	Purchase
25	25,000	Yes
35	50,000	Yes
32	35,000	No

Testing Data



Age	Income	Purchase
27	35,000	Yes
23	20,000	No
45	34,000	No



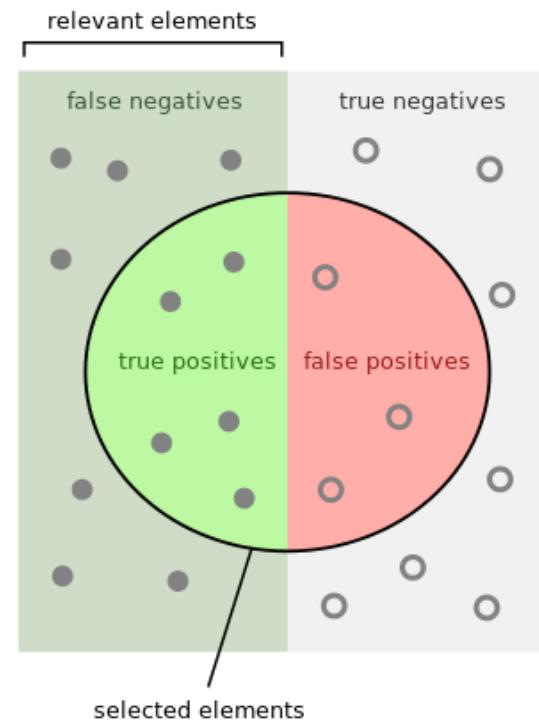


Confusion Matrix

Precision, Recall, F1

n=165	Predicted: NO	Predicted: YES	
Actual: NO	TN = 50	FP = 10	60
Actual: YES	FN = 5	TP = 100	105
	55	110	

- Precision = correctly predict = $TP / (TP + FP)$
- Recall = coverage = $TP / (TP + FN)$
- F1 = $(2 * \text{pre} * \text{rec}) / (\text{pre} + \text{rec})$



How many selected items are relevant?

$$\text{Precision} = \frac{\text{green}}{\text{green} + \text{red}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{green}}{\text{green} + \text{grey}}$$

+

Thank you & any questions